

# DATA 6320 - Scenario 3

## Program Sections

- Look for outliers
- Remove outliers
- Recode categorical data - label encoder
- Recode Categorical using a Defined Function
- Recode Categorical using dummy variables
- Recode Categorical using by merging two datasets

## Contents

- **Business Understanding**
  - Available resources, problems, goals
- **Data Understanding**
  - What data do you have available to you?
  - Install or Load tools or applications
    - Programming – Jupyter notebooks for Python or R Studio for R
    - BI/spreadsheets – Excel – PowerPivot - Tableau
  - Import or download the data
  - Format the data
  - View, explore, and summarize the data
- **Data Preparation**
  - Remove Columns
  - Remove Rows
  - Fill in null values
  - Replace or remove mistakes
  - **Remove outliers**
  - **Recode categorical or numerical features**
    - Construct new data feature engineering
    - For Supervised Learning, create X and Y
- **Modeling**
  - Split the data (Train/Test Split)
  - Transform the data
  - Setup models for machine learning/AI processes
  - Can also include developing the outline for visuals, dashboards or reports
- **Evaluation**
  - Hyper-parameter tuning
- **Deployment of models**

# BUSINESS UNDERSTANDING

---

## Business Objective

- What is the relationship between annual income and loan amount. Is it a good predictor?
- Can we predict the amount for a loan?
- What features are most important in predicting loan amount?

## Technical Objective

- Review different data cleansing techniques to continue to improve
  - Conduct a simple regression using scikit-learn
  - Conduct a multiple regression using statsmodels
  - Learn and perform Lasso and Ridge regression and tune its parameters
- 

# DATA UNDERSTANDING

---

## Import Libraries

In [1]:

```
#Code Block 01
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

pd.set_option('display.max_columns',500)

plt.style.use('seaborn-v0_8-colorblind') #a style that can be used for plots
```

```
sns.set_style('whitegrid')

In [2]:
#obtain the list of available styles

plt.style.available
```

```
Out[2]:
['Solarize_Light2',
 '_classic_test_patch',
 '_mpl-gallery',
 '_mpl-gallery-nogrid',
 'bmh',
 'classic',
 'dark_background',
 'fast',
 'fivethirtyeight',
 'ggplot',
 'grayscale',
 'seaborn-v0_8',
 'seaborn-v0_8-bright',
 'seaborn-v0_8-colorblind',
 'seaborn-v0_8-dark',
 'seaborn-v0_8-dark-palette',
 'seaborn-v0_8-darkgrid',
 'seaborn-v0_8-deep',
 'seaborn-v0_8-muted',
 'seaborn-v0_8-notebook',
 'seaborn-v0_8-paper',
 'seaborn-v0_8-pastel',
 'seaborn-v0_8-poster',
 'seaborn-v0_8-talk',
 'seaborn-v0_8-ticks',
 'seaborn-v0_8-white',
 'seaborn-v0_8-whitegrid',
 'tableau-colorblind10']
```

Import the data and create dataframes that can then be cleansed, recoded, transformed, and split.

## Import Data

---

[Top](#)

## Look for outliers

```
In [3]:
#Code Block 02

df_loandata_clean = pd.read_csv('data/DATA6320_Scenario3.csv', index_col = 0, header = 0
df_loandata_clean.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 37132 entries, 0 to 37131
Data columns (total 27 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   member_id        37132 non-null  int64   
 1   loan_amnt       37132 non-null  int64   
 2   term             37132 non-null  int64   
 3   int_rate         37132 non-null  float64  
 4   annual_inc       37132 non-null  float64  
 5   delinq_2yrs      37132 non-null  int64   
 6   inq_last_6mths   37132 non-null  int64   
 7   mths_since_last_delinq 37132 non-null  int64   
 8   mths_since_last_record 37132 non-null  int64   
 9   open_acc          37132 non-null  int64   
 10  pub_rec            37132 non-null  int64   
 11  revol_bal         37132 non-null  int64   
 12  total_acc         37132 non-null  int64   
 13  total_debt_paid   37132 non-null  float64  
 14  princ_int_ratio   37132 non-null  float64  
 15  collections_12_mths_ex_med 37132 non-null  int64   
 16  mths_since_last_major_derog 37132 non-null  int64   
 17  acc_now_delinq     37132 non-null  int64   
 18  tot_coll_amt       37132 non-null  int64   
 19  tot_cur_bal        37132 non-null  int64   
 20  total_credit_rv    37132 non-null  int64   
 21  revol_util         37132 non-null  float64  
 22  sub_grade          37132 non-null  object  
 23  emp_length          37132 non-null  int64   
 24  home_ownership      37132 non-null  object  
 25  loan_status         37132 non-null  object  
 26  reason              37132 non-null  object  
dtypes: float64(5), int64(18), object(4)
memory usage: 7.9+ MB

```

In [4]:

#Code Block 03

```

plt.figure(figsize=(20,16))

plt.subplot(221)
plt.title('Annual Income', fontweight='bold', color = 'green', fontsize='17', horizontalalignment='center')
sns.histplot(df_loandata_clean['annual_inc'], color="g", bins = 20, kde=True)

plt.subplot(222)
plt.title('Revolving Balance', fontweight='bold', color = 'blue', fontsize='17', horizontalalignment='center')
sns.histplot(df_loandata_clean['revol_bal'], color="b", bins = 20, kde=True)

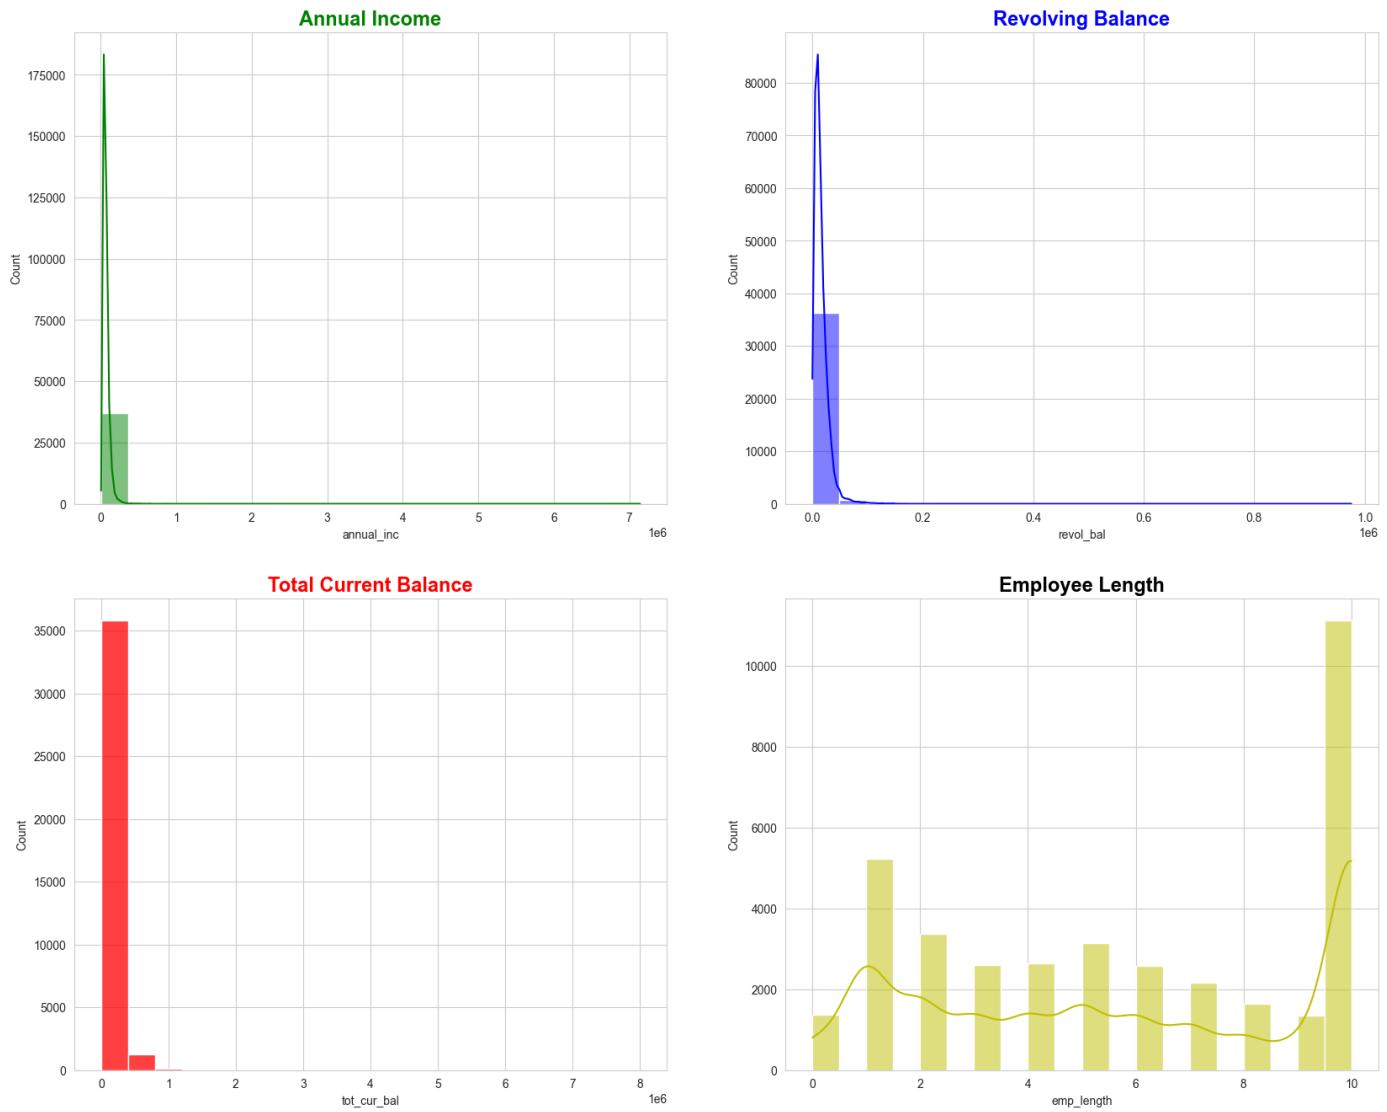
plt.subplot(223)
plt.title('Total Current Balance', fontweight='bold', color = 'red', fontsize='17', horizontalalignment='center')
sns.histplot(df_loandata_clean['tot_cur_bal'], color="r", bins = 20, kde=False)

plt.subplot(224)
plt.title('Employee Length', fontweight='bold', color = 'black', fontsize='17', horizontalalignment='center')
sns.histplot(df_loandata_clean['emp_length'], color="y", bins = 20, kde=True )

```

Out[4]:

```
<Axes: title={'center': 'Employee Length'}, xlabel='emp_length', ylabel='Count'>
```



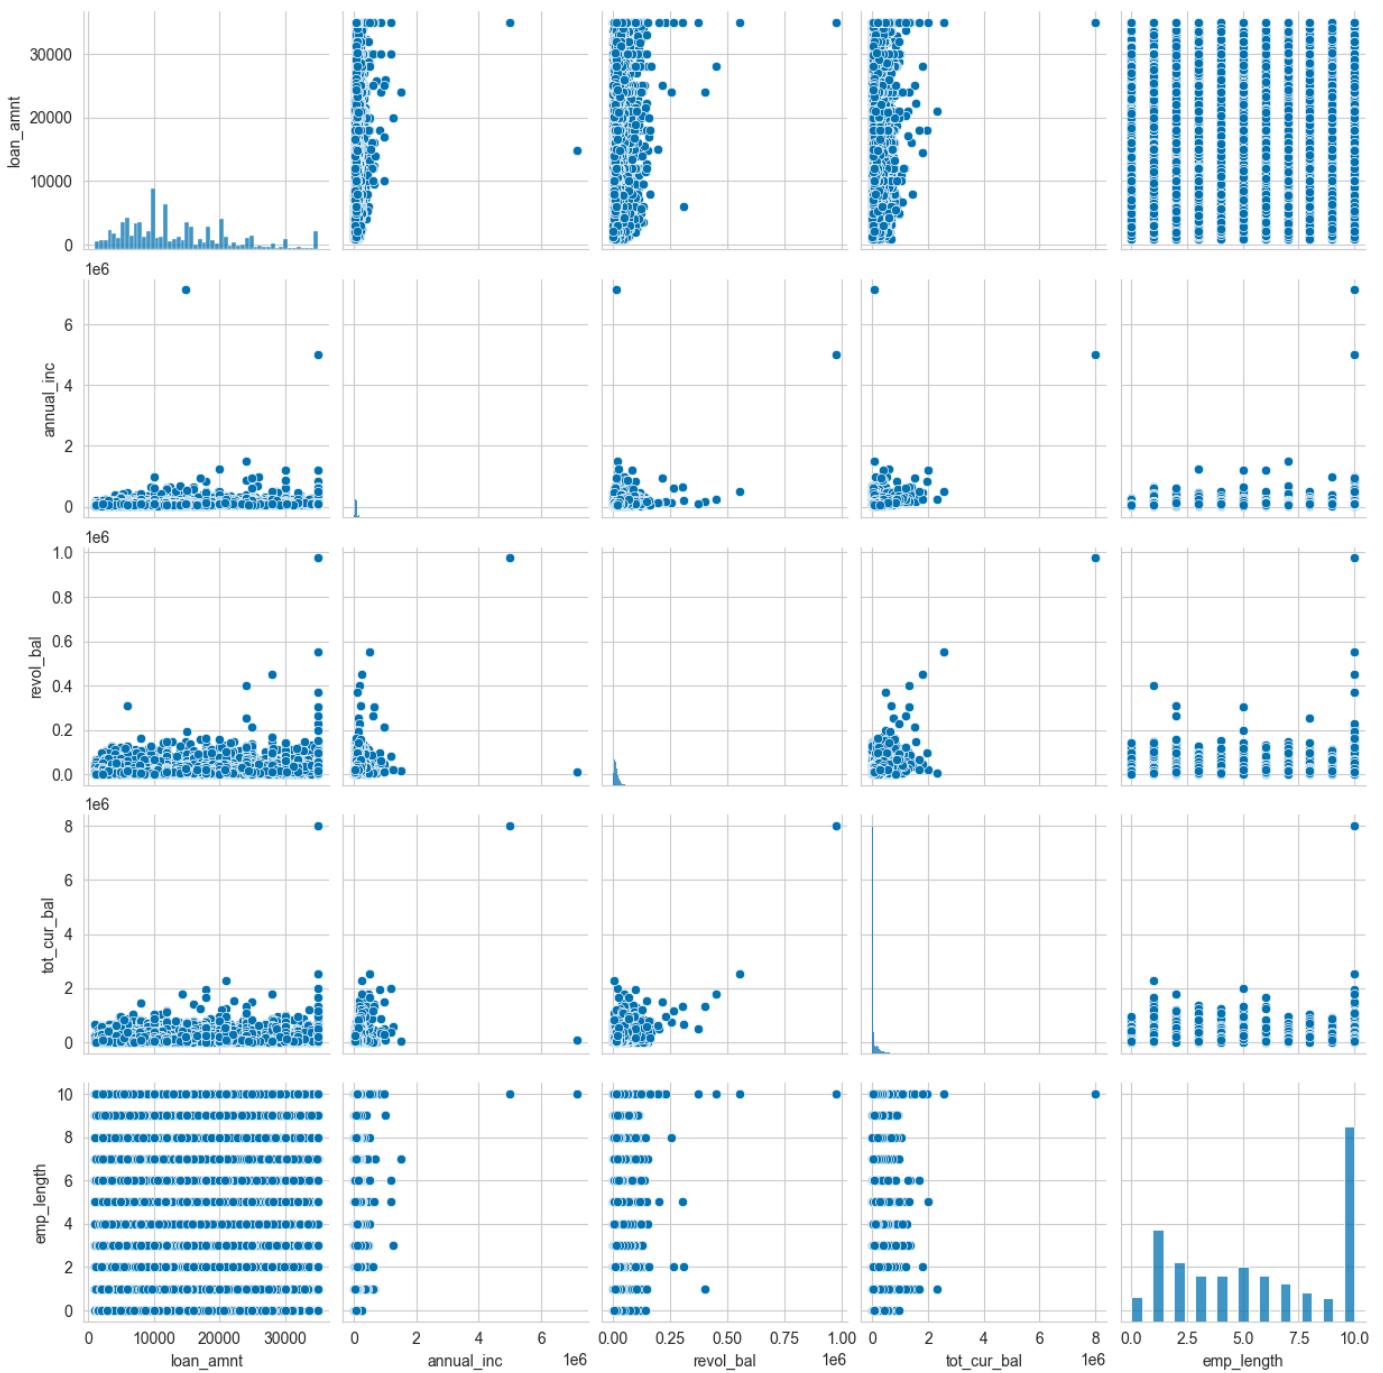
In [5]:

#Code Block 04

```
df_outliers = df_loandata_clean[['loan_amnt', 'annual_inc', 'revol_bal', 'tot_cur_bal', 'emp_length']]
sns.pairplot(df_outliers)
```

Out[5]:

<seaborn.axisgrid.PairGrid at 0x7fedf8197280>



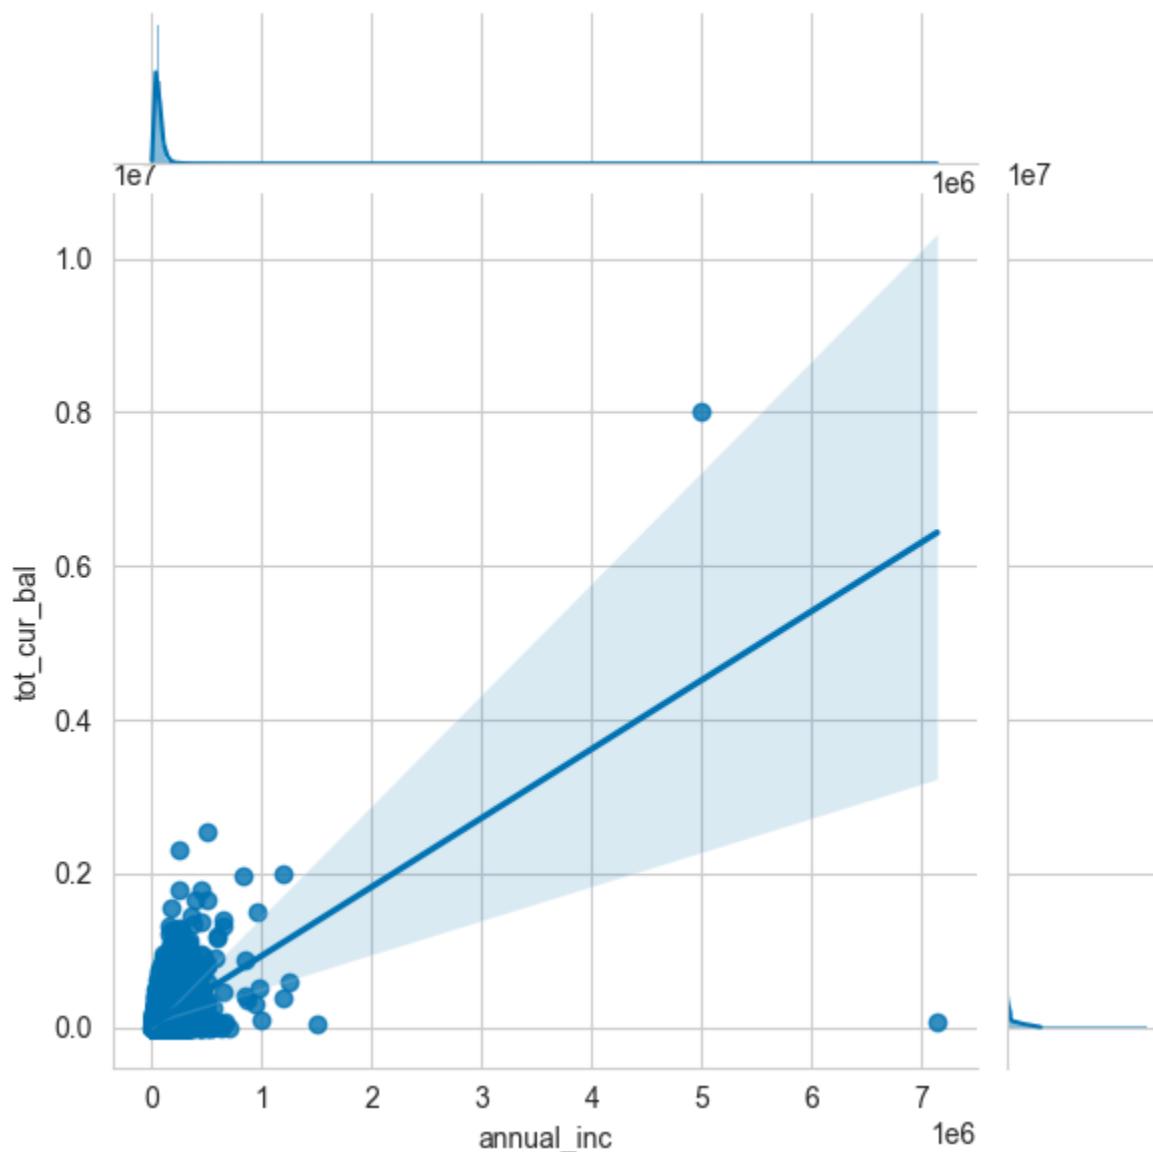
In [6]:

#Code Block 05

```
sns.jointplot(x="annual_inc", y="tot_cur_bal", data=df_loandata_clean, kind='reg')
```

Out[6]:

<seaborn.axisgrid.JointGrid at 0x7feea1be380>



In [7]:

#Code Block 06

```
df_loandata_clean['annual_inc'].value_counts()
```

Out[7]:

```
annual_inc
60000.0    1456
50000.0    1395
40000.0    1165
45000.0    1137
65000.0    1110
...
79445.0      1
92200.0      1
70584.0      1
36717.0      1
39990.0      1
Name: count, Length: 4043, dtype: int64
```

In [8]:

#Code Block 07

```
df_annual_inc = df_loandata_clean['annual_inc'].value_counts()
```

```
df_annual_inc = pd.DataFrame(df_annual_inc).reset_index()
print(df_annual_inc.columns)
#df_annual_inc.sort_values(by='index', ascending=False).head(5)
df_annual_inc.sort_values(by='annual_inc', ascending=False).head(5)
```

Index(['annual\_inc', 'count'], dtype='object')

Out[8]:

	annual_inc	count
1913	7141778.0	1
2031	5000000.0	1
3609	1500000.0	1
2413	1250000.0	1
698	1200000.0	2

**Note:** In the lecture, we used 'index' to sort values after calling reset\_index(), as in:

```
df_annual_inc = df_loandata_clean['annual_inc'].value_counts() df_annual_inc =
pd.DataFrame(df_annual_inc).reset_index() df_annual_inc.sort_values(by='index',
ascending=False).head(5)
```

However, depending on the version of pandas or the structure of the resulting DataFrame, the default column name 'index' may not be preserved. For example, after reset\_index(), the columns are named something like ['annual\_inc', 'count'] instead, as you see above.

To avoid errors, always check the column names using df.columns, and sort using the actual column name (e.g., 'annual\_inc') instead of 'index'.

Same output with one line of code

In [9]:

#Code Block 08

```
pd.DataFrame(df_loandata_clean['annual_inc'].value_counts()).reset_index().sort_values(b
```

Out[9]:

	annual_inc	count
1913	7141778.0	1
2031	5000000.0	1
3609	1500000.0	1
2413	1250000.0	1
698	1200000.0	2

Look at the two annual incomes that are > 1000000.0

- Is it worth keeping these six data points?

In [17]:

```
#Code Block 09
```

```
df_loandata_clean[df_loandata_clean['annual_inc'] >=1000000]
```

Out[17]:

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last
9523	1797692	35000	36	12.12	1200000.0	0	2	
18299	1407521	24000	60	21.00	1500000.0	0	0	
30802	2432841	20000	36	8.90	1250000.0	0	2	
35284	2839463	35000	36	15.31	5000000.0	1	2	
36011	2926819	30000	36	11.14	1200000.0	0	0	
36944	3267055	14825	36	13.11	7141778.0	0	2	

In [28]:

```
#Code Block 10
```

```
pd.DataFrame(df_loandata_clean['revol_bal'].value_counts()).reset_index().sort_values(by
```

Out[28]:

	revol_bal	count
13028	975800	1
12409	552758	1
11456	451481	1
11165	401258	1
12433	371817	1
12473	308071	1
12307	303993	1
12446	264260	1
11143	252191	1
12771	229112	1

Look at the two revol\_bal that are > 200000

- Is it worth keeping data points above 200,000?

In [29]:

```
#Code Block 11
```

```
df_loandata_clean[df_loandata_clean['revol_bal'] >=200000]
```

Out[29]:

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last
32964	2730966	25000	36	7.62	954285.00	0	0	

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last
33684	2739932	24000	36	6.62	160000.00	0	3	
33749	2740605	24000	60	17.27	141000.00	4	1	
33906	2742178	28000	36	17.77	250000.00	0	5	
34510	2829819	35000	36	18.75	175000.00	0	2	
35284	2839463	35000	36	15.31	5000000.00	1	2	
35982	2917830	35000	60	16.29	650000.00	0	1	
36558	3017184	6000	36	16.29	200000.00	0	0	
36940	3266877	35000	60	16.29	92974.96	0	0	
36982	3376772	35000	36	11.14	600000.00	0	1	
37077	3417245	35000	60	16.29	500000.00	0	1	

In [31]:

```
#Code Block 12
```

```
pd.DataFrame(df_loandata_clean['tot_cur_bal'].value_counts()).reset_index().sort_values(
```

Out[31]:

	tot_cur_bal	count
7720	8000078	1
8941	2547166	1
19902	2302431	1
18070	2008009	1
18260	1967533	1
11751	1787296	1
12688	1785763	1
10901	1662481	1
7713	1657520	1
5236	1549458	1

In [32]:

```
df_loandata_clean.head()
```

Out[32]:

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last_deli
0	1581986	9000	36	12.12	45000.0	0	3	
1	1751708	6625	36	11.14	28000.0	1	0	
2	1666916	9800	36	12.12	50000.0	0	0	
3	1758003	4250	36	8.90	38000.0	2	3	
4	1730191	16000	36	7.90	60000.0	0	0	

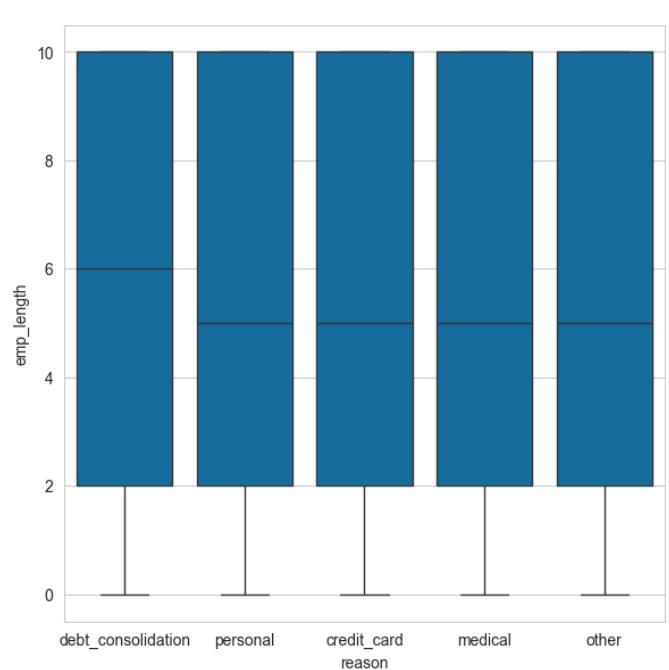
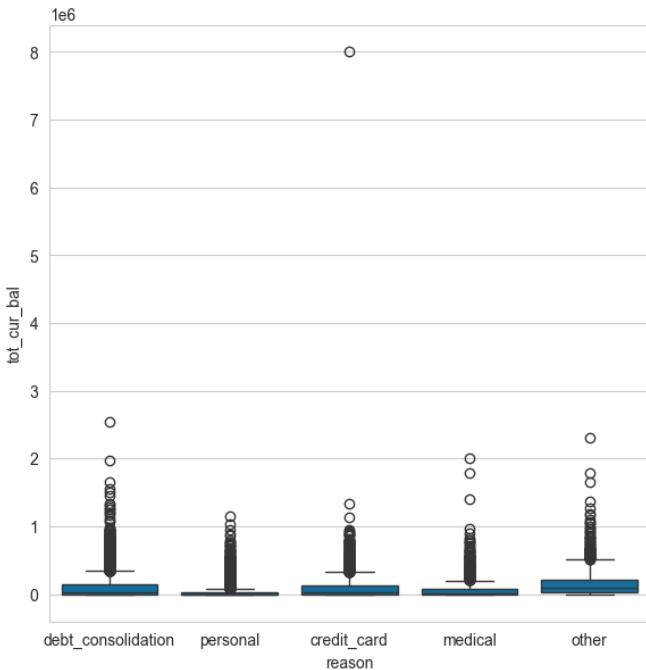
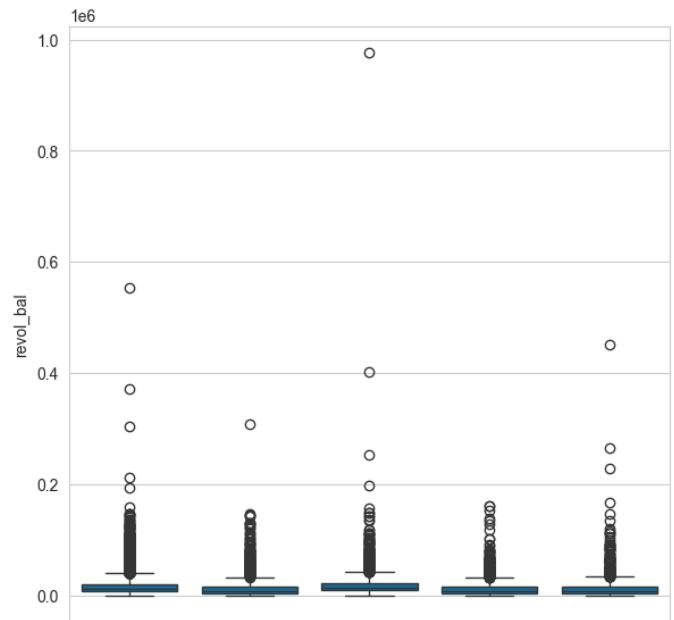
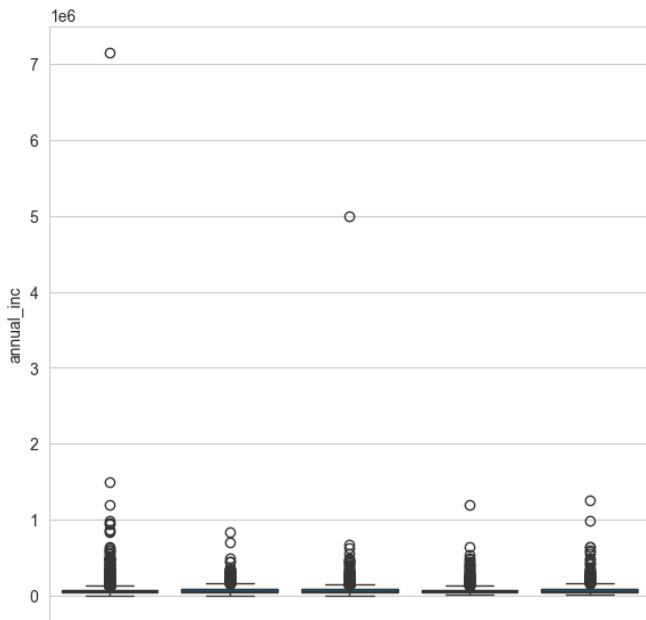
In [33]:

#Code Block 13

```
f, axes = plt.subplots(2, 2, figsize=(15, 15), sharex=True)
sns.boxplot(y='annual_inc', x = 'reason', data=df_loandata_clean, ax=axes[0,0])
sns.boxplot(y='revol_bal', x = 'reason', data=df_loandata_clean, ax=axes[0,1])
sns.boxplot(y='tot_cur_bal', x = 'reason', data=df_loandata_clean, ax=axes[1,0])
sns.boxplot(y='emp_length', x = 'reason', data=df_loandata_clean, ax=axes[1,1])
```

Out[33]:

```
<Axes: xlabel='reason', ylabel='emp_length'>
```



Top

## Remove outliers

Why are we removing observations?

- Appleton gives micro loans to individuals that need help with debt consolidation, medical bills, etc. They are normal people that just need a short term loan. Therefore, the objective is not to predict outliers, but instead the average person that asks for a loan. Outliers may skew the results.

Therefore:

- Remove records with annual\_inc > 1,000,000
- Remove records with revol\_bal > 200,000
- Remove records with total current balance > 2,000,000
- For employee length we change that to categories and then dummy variables

In [34]:

```
#Code Block 14
```

```
df_loandata_clean.shape
```

Out[34]:

(37132, 27)

In [35]:

```
#Code Block 15
```

```
df_loandata_clean.drop(df_loandata_clean[df_loandata_clean.annual_inc > 1000000].index,
df_loandata_clean.drop(df_loandata_clean[df_loandata_clean.revol_bal > 200000].index,
#df_loandata_clean.drop(df_loandata_clean[df_loandata_clean.tot_cur_bal > 2000000].index
df_loandata_clean.shape
```

Out[35]:

(37116, 27)

In [36]:

```
#Code Block 16
```

```
df_loandata_clean = df_loandata_clean[df_loandata_clean['tot_cur_bal'] < 2000000]
df_loandata_clean.shape
```

Out[36]:

(37115, 27)

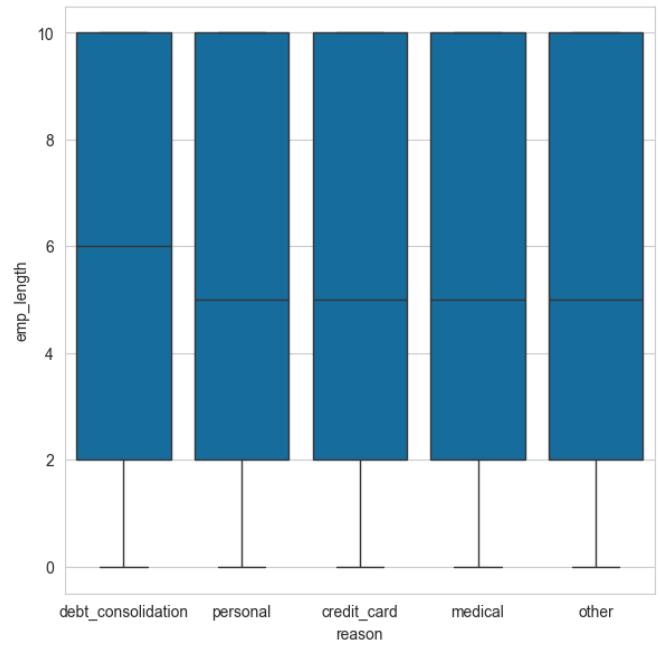
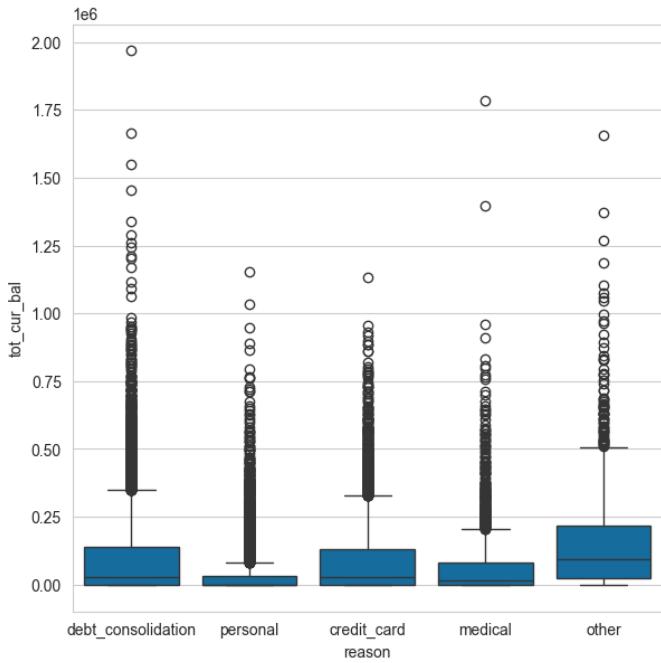
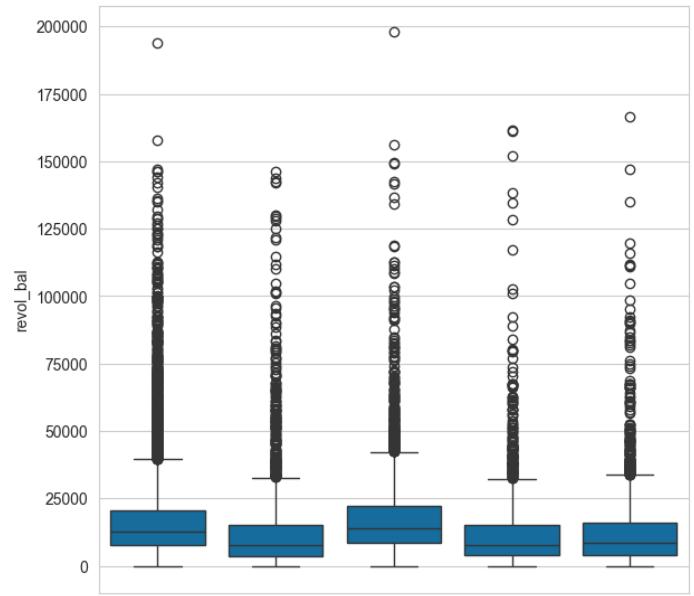
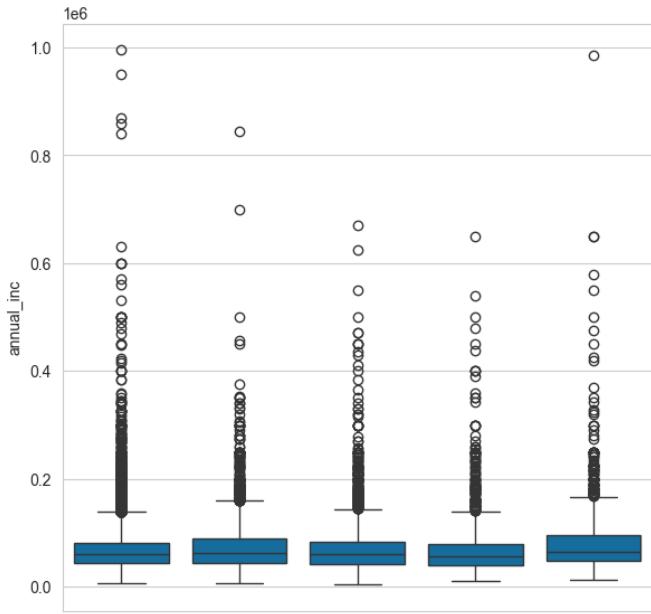
In [37]:

```
#Code Block 17
```

```
f, axes = plt.subplots(2, 2, figsize=(15, 15), sharex=True)
sns.boxplot(y='annual_inc', x = 'reason', data=df_loandata_clean, ax=axes[0,0])
sns.boxplot(y='revol_bal', x = 'reason', data=df_loandata_clean, ax=axes[0,1])
sns.boxplot(y='tot_cur_bal', x = 'reason', data=df_loandata_clean, ax=axes[1,0])
sns.boxplot(y='emp_length', x = 'reason', data=df_loandata_clean, ax=axes[1,1])
```

Out[37]:

<Axes: xlabel='reason', ylabel='emp\_length'>



## Re-code the Data

Top

### Recode categorical data - label encoder

#### 4 ways to recode categorical data

- **label encoder** - changes categories to integers based on alphabetical order

- **hot one encoder** - changes one column of categorical data into several binary (dummy) columns
- **use a custom function** for changing categories to integers
- **merge** from a separate dataset

In [38]:

#Code Block 18

```
df_loandata_clean.info()

<class 'pandas.core.frame.DataFrame'>
Index: 37115 entries, 0 to 37131
Data columns (total 27 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   member_id        37115 non-null   int64  
 1   loan_amnt        37115 non-null   int64  
 2   term              37115 non-null   int64  
 3   int_rate          37115 non-null   float64 
 4   annual_inc        37115 non-null   float64 
 5   delinq_2yrs       37115 non-null   int64  
 6   inq_last_6mths   37115 non-null   int64  
 7   mths_since_last_delinq 37115 non-null   int64  
 8   mths_since_last_record 37115 non-null   int64  
 9   open_acc          37115 non-null   int64  
 10  pub_rec            37115 non-null   int64  
 11  revol_bal         37115 non-null   int64  
 12  total_acc         37115 non-null   int64  
 13  total_debt_paid   37115 non-null   float64 
 14  princ_int_ratio   37115 non-null   float64 
 15  collections_12_mths_ex_med 37115 non-null   int64  
 16  mths_since_last_major_derog 37115 non-null   int64  
 17  acc_now_delinq    37115 non-null   int64  
 18  tot_coll_amt      37115 non-null   int64  
 19  tot_cur_bal       37115 non-null   int64  
 20  total_credit_rv   37115 non-null   int64  
 21  revol_util        37115 non-null   float64 
 22  sub_grade          37115 non-null   object  
 23  emp_length         37115 non-null   int64  
 24  home_ownership     37115 non-null   object  
 25  loan_status         37115 non-null   object  
 26  reason             37115 non-null   object  
dtypes: float64(5), int64(18), object(4)
memory usage: 7.9+ MB
```

## Convert all object features:

- sub\_grade - label encoder
- loan\_status - merge with new dataset
- reason - one-hot encoding
- home\_ownership - one-hot encoding

## Convert non-linear features

- emp\_length - will be converted to an object with a defined function and then encode with one-hot encoding

## Recode using Label Encoding

- Change subgrade to a numerical value and assume it is linear

In [39]:

```
#Code Block 19

from sklearn.preprocessing import LabelEncoder
lc = LabelEncoder()
```

In [40]:

```
#Code Block 20

df_loandata_clean.sub_grade.value_counts()
```

Out[40]:

```
sub_grade
B3      3766
B4      2838
C1      2529
B5      2494
C2      2370
B2      2322
A4      1885
A5      1696
B1      1651
C4      1377
D1      1317
D2      1294
C3      1279
C5      1151
D3      1050
A3      1010
A1      1001
D4      991
A2      896
D5      803
E1      558
E2      522
E4      427
E3      418
E5      336
```

```
F1      286  
F2      232  
F3      175  
F4      131  
F5      124  
G1      86  
G2      51  
G3      23  
G4      17  
G5       9  
Name: count, dtype: int64
```

In [41]:

```
#Code Block 21
```

```
df_loandata_clean.head()
```

Out[41]:

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last_deli
0	1581986	9000	36	12.12	45000.0	0	3	
1	1751708	6625	36	11.14	28000.0	1	0	
2	1666916	9800	36	12.12	50000.0	0	0	
3	1758003	4250	36	8.90	38000.0	2	3	
4	1730191	16000	36	7.90	60000.0	0	0	

In [42]:

```
#Code Block 22
```

```
df_loandata_clean['sub_grade'] = lc.fit_transform(df_loandata_clean['sub_grade'])
```

In [43]:

```
#Code Block 23
```

```
df_loandata_clean['sub_grade'].value_counts()
```

Out[43]:

```
sub_grade  
7      3766  
8      2838  
10     2529  
9      2494  
11     2370  
6      2322  
3      1885  
4      1696  
5      1651  
13     1377  
15     1317  
16     1294  
12     1279  
14     1151  
17     1050  
2      1010  
0      1001  
18     991
```

```
1      896  
19     803  
20     558  
21     522  
23     427  
22     418  
24     336  
25     286  
26     232  
27     175  
28     131  
29     124  
30      86  
31      51  
32      23  
33      17  
34       9  
Name: count, dtype: int64
```

---

Top

## Recode Categorical using a Defined Function

In [44]:

#Code Block 24

```
df_loandata_clean['emp_length'].value_counts()
```

Out[44]:

```
emp_length  
10     11104  
1      5219  
2      3359  
5      3127  
4      2632  
3      2590  
6      2573  
7      2165  
8      1635  
0      1375  
9      1336  
Name: count, dtype: int64
```

### How to recode?

- 10 refers to 10 or more
- By talking to management, they believe that this can be changed to:
  - **10 or more:** 10
  - **7 to 9:** 7, 8, and 9
  - **4 to 6:** 4, 5, and 6
  - **3 or less:** 0, 1, 2, and 3

In [45]:

```
#Code Block 25
```

```
def EmpLength(d):
    if d['emp_length'] == 10:
        return "10 or more"
    elif d['emp_length'] >= 7:
        return "7 to 9"
    elif d['emp_length'] >= 4:
        return "4 to 6"
    else:
        return "3 or less"
df_loandata_clean['emp_length_cat'] = df_loandata_clean.apply(EmpLength, axis = 1)
df_loandata_clean[['emp_length', 'emp_length_cat']].head()
```

```
Out[45]:
```

	emp_length	emp_length_cat
0	6	4 to 6
1	3	3 or less
2	10	10 or more
3	4	4 to 6
4	10	10 or more

```
In [46]:
```

```
df_loandata_clean.head()
```

```
Out[46]:
```

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last_deli
0	1581986	9000	36	12.12	45000.0	0	3	
1	1751708	6625	36	11.14	28000.0	1	0	
2	1666916	9800	36	12.12	50000.0	0	0	
3	1758003	4250	36	8.90	38000.0	2	3	
4	1730191	16000	36	7.90	60000.0	0	0	

---

```
Top
```

## Recode Categorical using dummy variables

How to automate the dropping of the largest dummy variable

```
In [47]:
```

```
#Code Block 26
```

```
df_loandata_clean['home_ownership'].value_counts()
```

```
Out[47]:
```

```
home_ownership
MORTGAGE    18081
RENT        16096
OWN         2938
Name: count, dtype: int64
```

In [48]:

#Code Block 27

```
df_dummy = pd.DataFrame(df_loandata_clean['home_ownership'].value_counts().reset_index())
df_dummy
```

Out[48]:

	home_ownership	count
0	MORTGAGE	18081
1	RENT	16096
2	OWN	2938

In [49]:

#Code Block 28

```
var_dummy = df_dummy.iloc[0, 0]
var_dummy
```

Out[49]:

'MORTGAGE'

In [50]:

#Code Block 29

```
var_dummy_prefix = "home"
var_dumpre = var_dummy_prefix + "_" + var_dummy
var_dumpre
```

Out[50]:

'home\_MORTGAGE'

In [53]:

#Code Block 29

```
#Added dtype=float to convert to display 0/1 values than the boolean
dummies = pd.get_dummies(df_loandata_clean['home_ownership'], drop_first = False, prefix
dummies.head()
```

Out[53]:

	home_MORTGAGE	home_own	home_RENT
0	0.0	0.0	1.0
1	0.0	0.0	1.0
2	0.0	0.0	1.0
3	0.0	1.0	0.0
4	1.0	0.0	0.0

In [54]:

#Code Block 30

```
dummies = dummies.drop(var_dumpre, axis = 1)
dummies.head()
```

Out[54]:

	home_OWNER	home_RENT
0	0.0	1.0
1	0.0	1.0
2	0.0	1.0
3	1.0	0.0
4	0.0	0.0

In [55]:

#Code Block 32

```
df_loandata_clean = pd.concat([df_loandata_clean, dummies], axis = 1)
df_loandata_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 37115 entries, 0 to 37131
Data columns (total 30 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   member_id        37115 non-null   int64  
 1   loan_amnt        37115 non-null   int64  
 2   term              37115 non-null   int64  
 3   int_rate          37115 non-null   float64 
 4   annual_inc        37115 non-null   float64 
 5   delinq_2yrs       37115 non-null   int64  
 6   inq_last_6mths   37115 non-null   int64  
 7   mths_since_last_delinq  37115 non-null   int64  
 8   mths_since_last_record 37115 non-null   int64  
 9   open_acc          37115 non-null   int64  
 10  pub_rec           37115 non-null   int64  
 11  revol_bal         37115 non-null   int64  
 12  total_acc         37115 non-null   int64  
 13  total_debt_paid   37115 non-null   float64 
 14  princ_int_ratio   37115 non-null   float64 
 15  collections_12_mths_ex_med 37115 non-null   int64  
 16  mths_since_last_major_derog 37115 non-null   int64  
 17  acc_now_delinq    37115 non-null   int64  
 18  tot_coll_amt     37115 non-null   int64  
 19  tot_cur_bal      37115 non-null   int64  
 20  total_credit_rv  37115 non-null   int64  
 21  revol_util       37115 non-null   float64 
 22  sub_grade         37115 non-null   int64  
 23  emp_length        37115 non-null   int64  
 24  home_ownership    37115 non-null   object  
 25  loan_status        37115 non-null   object  
 26  reason             37115 non-null   object  
 27  emp_length_cat    37115 non-null   object  
 28  home_OWNER         37115 non-null   float64 
 29  home_RENT          37115 non-null   float64 
```

```
dtypes: float64(7), int64(19), object(4)
memory usage: 8.8+ MB
```

## Create dummy variables in one code block

In [57]:

```
#Code Block 33

#Create prefix for dummy variables
var_dummy_prefix = "reason"

#Find the highest count category
df_dummy = pd.DataFrame(df_loandata_clean['reason'].value_counts().reset_index())
var_dummy = df_dummy.iloc[0, 0]

#Create variable to drop the highest count column
var_dumpre = var_dummy_prefix + "_" + var_dummy

#Create dummy variables
dummies = pd.get_dummies(df_loandata_clean['reason'], drop_first = False, prefix=var_dummy)

#Drop the highest count dummy variable
dummies = dummies.drop(var_dumpre, axis = 1)

#Concat the dummy variables to the main dataset
df_loandata_clean = pd.concat([df_loandata_clean, dummies], axis = 1)

#Display the variables and new dataset

display(df_dummy)
print("-----")
print("Highest Count:")
print(var_dummy)
print("-----")

display(dummies.head())
df_loandata_clean.head()
```

	reason	count
0	debt_consolidation	22312
1	credit_card	7234
2	personal	3121
3	other	2270
4	medical	2178

-----  
Highest Count:  
debt\_consolidation  
-----

	reason_credit_card	reason_medical	reason_other	reason_personal
0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	1.0

	reason_credit_card	reason_medical	reason_other	reason_personal
2	1.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	1.0

Out[57]:

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last_deli
0	1581986	9000	36	12.12	45000.0	0	3	
1	1751708	6625	36	11.14	28000.0	1	0	
2	1666916	9800	36	12.12	50000.0	0	0	
3	1758003	4250	36	8.90	38000.0	2	3	
4	1730191	16000	36	7.90	60000.0	0	0	

In [58]:

#Code Block 34

```
#Create prefix for dummy variables
var_dummy_prefix = "emp"

#Find the highest count category
df_dummy = pd.DataFrame(df_loandata_clean['emp_length_cat'].value_counts().reset_index())
var_dummy = df_dummy.iloc[0, 0]

#Create variable to drop the highest count column
var_dumpre = var_dummy_prefix + "_" + var_dummy

#Create dummy variables
dummies = pd.get_dummies(df_loandata_clean['emp_length_cat'], drop_first = False, prefix = var_dumpre)

#Drop the highest count dummy variable
dummies = dummies.drop(var_dumpre, axis = 1)

#Concat the dummy variables to the main dataset
df_loandata_clean = pd.concat([df_loandata_clean, dummies], axis = 1)

#Display the variables and new dataset
display(df_dummy)
print("-----")
print("Highest Count:")
print(var_dummy)
print("-----")

display(dummies.head())
df_loandata_clean.head()
```

	emp_length_cat	count
0	3 or less	12543
1	10 or more	11104

emp_length_cat	count	
2	4 to 6	8332
3	7 to 9	5136

Highest Count:

3 or less

	emp_10 or more	emp_4 to 6	emp_7 to 9
0	0.0	1.0	0.0
1	0.0	0.0	0.0
2	1.0	0.0	0.0
3	0.0	1.0	0.0
4	1.0	0.0	0.0

Out[58]:

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last_deli
0	1581986	9000	36	12.12	45000.0	0	3	
1	1751708	6625	36	11.14	28000.0	1	0	
2	1666916	9800	36	12.12	50000.0	0	0	
3	1758003	4250	36	8.90	38000.0	2	3	
4	1730191	16000	36	7.90	60000.0	0	0	

## What to do with Term?

- Term is the number of months for each loan (36 or 60 months)
- Since it is not a linear relationship, it is actually a label (based on numbers), it should be treated as a label/categorical variable.

In [59]:

#Code Block 35

```
df_loandata_clean['term'].value_counts()
```

Out[59]:

term	count
36	30325
60	6790
Name: count, dtype: int64	

## How to code?

- Since most loans are for 36 months, we will want to see if 60 month term loans have an effect on the loan amount.

In [61]:

### #Code Block 36

```
#Create prefix for dummy variables
var_dummy_prefix = "term"

#Find the highest count category
df_dummy = pd.DataFrame(df_loandata_clean['term'].value_counts().reset_index())
var_dummy = df_dummy.iloc[0, 0]

#Create variable to drop the highest count column

#NOTE: Since 60 is an integer, you need to set it as a string before you can concatenate
var_dumpre = var_dummy_prefix + "_" + str(var_dummy)

#Create dummy variables
dummies = pd.get_dummies(df_loandata_clean['term'], drop_first = False, prefix=var_dummy)

#Drop the highest count dummy variable
dummies = dummies.drop(var_dumpre, axis = 1)

#Concat the dummy variables to the main dataset
df_loandata_clean = pd.concat([df_loandata_clean, dummies], axis = 1)

#Display the variables and new dataset

display(df_dummy)
print("-----")
print("Highest Count:")
print(var_dummy)
print("-----")

display(dummies.head())
df_loandata_clean.head()
```

term	count
0	36 30325
1	60 6790

Highest Count:

36

term_60
0 0.0
1 0.0
2 0.0
3 0.0
4 0.0

Out[61]:

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last_deli
0	1581986	9000	36	12.12	45000.0	0	3	
1	1751708	6625	36	11.14	28000.0	1	0	
2	1666916	9800	36	12.12	50000.0	0	0	
3	1758003	4250	36	8.90	38000.0	2	3	
4	1730191	16000	36	7.90	60000.0	0	0	

Top

## Recode Categorical using by merging two datasets

In [62]:

```
#Code Block 37
#url = 'https://data63206330.file.core.windows.net/data6320/Loan_is_Bad.csv?sp=rl&st=202
df_loanisbad = pd.read_csv("data/Loan_is_Bad.csv", index_col = None, header = 0)
df_loanisbad
```

Out[62]:

	loan_status	loan_is_bad
0	Charged Off	1
1	Current	0
2	Fully Paid	0
3	In Grace Period	0
4	Late (16-30 days)	0
5	Late (31-120 days)	1
6	Default	1

In [63]:

```
#Code Block 38
df_loandata_clean['loan_status'].value_counts()
```

Out[63]:

```
loan_status
Fully Paid      28510
Charged Off     5502
Current         2775
Late (31-120 days)  213
In Grace Period   87
Late (16-30 days)  22
Default          6
Name: count, dtype: int64
```

In [64]:

### #Code Block 39

```
df_loandata_clean = pd.merge(df_loandata_clean, df_loanisbad, on='loan_status', how='lef
```

In [65]:

### #Code Block 40

```
df_loandata_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 37115 entries, 0 to 37114
Data columns (total 44 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   member_id        37115 non-null   int64  
 1   loan_amnt        37115 non-null   int64  
 2   term              37115 non-null   int64  
 3   int_rate          37115 non-null   float64 
 4   annual_inc        37115 non-null   float64 
 5   delinq_2yrs       37115 non-null   int64  
 6   inq_last_6mths   37115 non-null   int64  
 7   mths_since_last_delinq 37115 non-null   int64  
 8   mths_since_last_record 37115 non-null   int64  
 9   open_acc          37115 non-null   int64  
 10  pub_rec           37115 non-null   int64  
 11  revol_bal         37115 non-null   int64  
 12  total_acc         37115 non-null   int64  
 13  total_debt_paid  37115 non-null   float64 
 14  princ_int_ratio  37115 non-null   float64 
 15  collections_12_mths_ex_med 37115 non-null   int64  
 16  mths_since_last_major_derog 37115 non-null   int64  
 17  acc_now_delinq   37115 non-null   int64  
 18  tot_coll_amt    37115 non-null   int64  
 19  tot_cur_bal     37115 non-null   int64  
 20  total_credit_rv 37115 non-null   int64  
 21  revol_util      37115 non-null   float64 
 22  sub_grade        37115 non-null   int64  
 23  emp_length       37115 non-null   int64  
 24  home_ownership  37115 non-null   object  
 25  loan_status      37115 non-null   object  
 26  reason            37115 non-null   object  
 27  emp_length_cat  37115 non-null   object  
 28  home_ownN        37115 non-null   float64 
 29  home_RENT         37115 non-null   float64 
 30  reason_credit_card 37115 non-null   bool   
 31  reason_medical   37115 non-null   bool   
 32  reason_other     37115 non-null   bool   
 33  reason_personal  37115 non-null   bool   
 34  reason_credit_card 37115 non-null   float64 
 35  reason_medical   37115 non-null   float64 
 36  reason_other     37115 non-null   float64 
 37  reason_personal  37115 non-null   float64 
 38  emp_10 or more   37115 non-null   float64 
 39  emp_4 to 6        37115 non-null   float64 
 40  emp_7 to 9        37115 non-null   float64 
 41  term_60           37115 non-null   float64 
 42  term_60           37115 non-null   float64
```

```
43 loan_is_bad           37115 non-null int64
dtypes: bool(4), float64(16), int64(20), object(4)
memory usage: 11.5+ MB
```

In [66]:

```
df_loandata_clean.head()
```

Out[66]:

	member_id	loan_amnt	term	int_rate	annual_inc	delinq_2yrs	inq_last_6mths	mths_since_last_deli
0	1581986	9000	36	12.12	45000.0	0	3	
1	1751708	6625	36	11.14	28000.0	1	0	
2	1666916	9800	36	12.12	50000.0	0	0	
3	1758003	4250	36	8.90	38000.0	2	3	
4	1730191	16000	36	7.90	60000.0	0	0	

## The Final Dataset

In [67]:

```
#Code Block 41
```

```
df_loandata_clean = df_loandata_clean.drop(['term','emp_length','home_ownership', 'loan_
df_loandata_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 37115 entries, 0 to 37114
Data columns (total 38 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   member_id        37115 non-null   int64  
 1   loan_amnt       37115 non-null   int64  
 2   int_rate         37115 non-null   float64 
 3   annual_inc      37115 non-null   float64 
 4   delinq_2yrs     37115 non-null   int64  
 5   inq_last_6mths  37115 non-null   int64  
 6   mths_since_last_delinq  37115 non-null   int64  
 7   mths_since_last_record 37115 non-null   int64  
 8   open_acc         37115 non-null   int64  
 9   pub_rec          37115 non-null   int64  
 10  revol_bal        37115 non-null   int64  
 11  total_acc        37115 non-null   int64  
 12  total_debt_paid 37115 non-null   float64 
 13  princ_int_ratio 37115 non-null   float64 
 14  collections_12_mths_ex_med 37115 non-null   int64  
 15  mths_since_last_major_derog 37115 non-null   int64  
 16  acc_now_delinq   37115 non-null   int64  
 17  tot_coll_amt    37115 non-null   int64  
 18  tot_cur_bal     37115 non-null   int64  
 19  total_credit_rv 37115 non-null   int64  
 20  revol_util      37115 non-null   float64 
 21  sub_grade        37115 non-null   int64  
 22  home OWN        37115 non-null   float64 
 23  home RENT        37115 non-null   float64 
 24  reason_credit_card 37115 non-null   bool
```

```
25 reason_medical           37115 non-null  bool
26 reason_other              37115 non-null  bool
27 reason_personal            37115 non-null  bool
28 reason_credit_card         37115 non-null  float64
29 reason_medical             37115 non-null  float64
30 reason_other                37115 non-null  float64
31 reason_personal              37115 non-null  float64
32 emp_10 or more             37115 non-null  float64
33 emp_4 to 6                  37115 non-null  float64
34 emp_7 to 9                  37115 non-null  float64
35 term_60                      37115 non-null  float64
36 term_60                      37115 non-null  float64
37 loan_is_bad                  37115 non-null  int64
dtypes: bool(4), float64(16), int64(18)
memory usage: 9.8 MB
```

In [68]:

#Code Block 42

```
#df_loandata_clean.to_csv('data/Appleton_CleanData.csv')
```

In [ ]:

In [ ]: